Questions:

1. Can we have different approaches for report and submission? (There can be one that is better for memory while the other gives better results while being more resource intensive)

   `-> They would love to know all the approaches and experiments we performed in the presentation and report. However, they recommend submitting the approach that gives the best results.

2. Where should we run our models?

   -> They suggested using free tier Google Colab as it would reflect on the constraint environment. There is a limit of 4Gb RAM for inference. For training, they haven't specified yet, though we should expect building solutions that work in resource constraint environment.

3. What happens if the memory limit of 4Gb is exceeded in inference?

   -> There will be a penalty, the magnitude of which is yet to be decided.

4. Evaluation queries:

   -> Given a question and a set of paragraphs, our model will need to predict if there exists an answer, the paragraph, and the exact answer. The predicted answer will be compared with the ground truth answer to get the F1 score. [F1 score for Paragraph selection and F1 score for actual answer]

   Some queries that were raised:
   1. There could be instances where multiple paragraphs answer the question.
   2. There could be multiple correct answers to the question:
      Eg 24th January 2002 or simply 2002 [an example from the dataset]

5. Submission queries:

   -> Midterm: Report is mandatory.
   -> End term: We need to submit the code along with the report. It will consist of two rounds as stated in the PS. It will be best if we submit Google Colab notebooks.